

**VISVESVARAYA
TECHNOLOGICAL UNIVERSITY**
“JnanaSangama”, Belgaum -590014, Karnataka.



LAB REPORT
on

Machine Learning (23CS6PCMAL)

Submitted by

Varsha Prasanth (1BM22CS321)

in partial fulfillment for the award of the degree of

BACHELOR OF ENGINEERING
in
COMPUTER SCIENCE AND ENGINEERING



B.M.S. COLLEGE OF ENGINEERING
(Autonomous Institution under VTU)
BENGALURU-560019
Sep-2024 to Jan-2025

B.M.S. College of Engineering,
Bull Temple Road, Bangalore 560019
(Affiliated To Visvesvaraya Technological University, Belgaum)
Department of Computer Science and Engineering



CERTIFICATE

This is to certify that the Lab work entitled “Machine Learning (23CS6PCMAL)” carried out by **Varsha Prasanth (1BM22CS321)**, who is a bonafide student of **B.M.S. College of Engineering**. It is in partial fulfillment for the award of **Bachelor of Engineering in Computer Science and Engineering** of the Visvesvaraya Technological University, Belgaum. The Lab report has been approved as it satisfies the academic requirements in respect of an Machine Learning (23CS6PCMAL) work prescribed for the said degree.

Ramya K M
Assistant Professor
Department of CSE , BMSCE

Dr. Kavitha Sooda
Professor & HOD
Department of CSE , BMSCE

Index

Sl. No.	Date	Experiment Title	Page No.
1	21-2-2025	Write a python program to import and export data using Pandas library functions	4
2	3-3-2025	Demonstrate various data pre-processing techniques for a given dataset	8
3	10-3-2025	Implement Linear and Multi-Linear Regression algorithm using appropriate dataset	11
4	17-3-2025	Build Logistic Regression Model for a given dataset	15
5	24-3-2025	Use an appropriate data set for building the decision tree (ID3) and apply this knowledge to classify a new sample.	18
6	7-4-2025	Build KNN Classification model for a given dataset.	21
7	21-4-2025	Build Support vector machine model for a given dataset	24
8	5-5-2025	Implement Random forest ensemble method on a given dataset.	27
9	5-5-2025	Implement Boosting ensemble method on a given dataset.	29
10	12-5-2025	Build k-Means algorithm to cluster a set of data stored in a .CSV file.	31
11	12-5-2025	Implement Dimensionality reduction using Principal Component Analysis (PCA) method.	33

Github Link:

<https://github.com/Varsha-Pr/ML-Lab>

Program 1

Write a python program to import and export data using Pandas library functions

Code:

```
import pandas as pd
data = {
'Name': ['Alice', 'Bob', 'Charlie', 'David'],
'Age': [25, 30, 35, 40],
'City': ['New York', 'Los Angeles', 'Chicago', 'Houston']
}
df = pd.DataFrame(data)
print("Sample data:")
print(df.head())
from sklearn.datasets import load_iris
iris = load_iris()
df = pd.DataFrame(iris.data, columns=iris.feature_names)
df['target'] = iris.target
print("Sample data:")
print(df.head())
file_path = 'data.csv'
df = pd.read_csv(file_path)
print("Sample data:")
print(df.head())
print("\n")
file_path = 'mobiles-dataset-2025.csv'
df = pd.read_csv(file_path, encoding='latin-1') # or 'cp1252' or other suitable encoding
print("Sample data:")
print(df.head())
import pandas as pd

data = {
'USN': ['IS001', 'IS002', 'IS003', 'IS004', 'IS005'],
'Name': ['Alice', 'Bob', 'Charlie', 'David', 'Eve'],
'Marks': [25, 30, 35, 40, 45]
}

df = pd.DataFrame(data)
print("Sample data:")
print(df.head())
from sklearn.datasets import load_diabetes
iris = load_diabetes()
df = pd.DataFrame(iris.data, columns=iris.feature_names)

print("Sample data:")
print(df.head())
```

```

file_path = 'sample_sales_data.csv'
df = pd.read_csv(file_path)
print("Sample data:")
print(df.head())
print("\n")

df = pd.read_csv("/content/dataset-of-diabetes .csv",encoding='latin-1')
print("Sample data:")
print(df.head())
print("\n")

df=pd.read_csv('sample_sales_data.csv')
print("Sample data:")
print(df.head())

df.to_csv('output.csv',index=False)
print("Data saved to output.csv")
sales_df =pd.read_csv('sample_sales_data.csv')
print("Sample data:")
print(sales_df.head())
sales_by_region =sales_df.groupby('Region')['Sales'].sum()
print("\nTotal sales by region:")
print(sales_by_region)
best_selling_products
=sales_df.groupby('Product')['Quantity'].sum().sort_values(ascending=False) print("\nBest-selling
products by quantity:")
print(best_selling_products)
sales_by_region.to_csv('sales_by_region.csv')
best_selling_products.to_csv('best_selling_products.csv')
print("Data saved to sales_by_region.csv and best_selling_products.csv")

import yfinance as yf
import matplotlib.pyplot as plt
tickers = ["RELIANCE.NS", "TCS.NS", "INFY.NS"]
data = yf.download(tickers, start="2022-10-01", end="2023-10-01",
group_by='ticker')
print("First 5 rows of the dataset:")
print(data.head())
print("\nShape of the dataset:")
print(data.shape)
print("\nColumn names:")
print(data.columns)
print("\n")
reliance_data = data['RELIANCE.NS']
print("\nSummary statistics for Reliance Industries:")
print(reliance_data.describe())
reliance_data['Daily Return'] = reliance_data['Close'].pct_change()
print("\n")
plt.figure(figsize=(12, 6))
plt.subplot(2, 1, 1)
reliance_data['Close'].plot(title="Reliance Industries - Closing Price")

```

```

plt.subplot(2, 1, 2)
reliance_data['Daily Return'].plot(title="Reliance Industries - Daily Returns", color='orange')
plt.tight_layout()
plt.show()
reliance_data.to_csv('reliance_stock_data.csv')

```

```

tickers = ["HDFCBANK.NS", "ICICI.NS", "KOTAKBANK.NS"]
data = yf.download(tickers, start="2024-01-01", end="2024-12-30",
group_by='ticker')
print("First 5 rows of the dataset:")
print(data.head())
print("\nShape of the dataset:")
print(data.shape)
print("\nColumn names:")
print(data.columns)
print("\n")
reliance_data = data['HDFCBANK.NS']
print("\nSummary statistics for Reliance Industries:")
print(reliance_data.describe())
reliance_data['Daily Return'] = reliance_data['Close'].pct_change()
print("\n")
plt.figure(figsize=(12, 6))
plt.subplot(2, 1, 1)
reliance_data['Close'].plot(title="HDFC Industries - Closing Price")
plt.subplot(2, 1, 2)
reliance_data['Daily Return'].plot(title="HDFC Industries - Daily Returns", color='red')
plt.tight_layout()
plt.show()
reliance_data.to_csv('hdfc_stock_data.csv')
print("\nhdfc stock data saved to 'hdfc_stock_data.csv'.")

```

```

tickers = ["HDFCBANK.NS", "ICICIBANK.NS", "KOTAKBANK.NS"]
data = yf.download(tickers, start="2024-01-01", end="2024-12-30",
group_by='ticker')
print("First 5 rows of the dataset:")
print(data.head())
print("\nShape of the dataset:")
print(data.shape)
print("\nColumn names:")
print(data.columns)
print("\n")
reliance_data = data['ICICIBANK.NS']
print("\nSummary statistics for ICICI Industries:")
print(reliance_data.describe())
reliance_data['Daily Return'] = reliance_data['Close'].pct_change()
print("\n")
plt.figure(figsize=(12, 6))
plt.subplot(2, 1, 1)
reliance_data['Close'].plot(title="ICICI Industries - Closing Price")
plt.subplot(2, 1, 2)
reliance_data['Daily Return'].plot(title="ICICI Industries - Daily Returns", color='BLACK')
plt.tight_layout()

```

```

plt.show()
reliance_data.to_csv('icici_stock_data.csv')
print("\nicici stock data saved to 'icici_stock_data.csv'.")

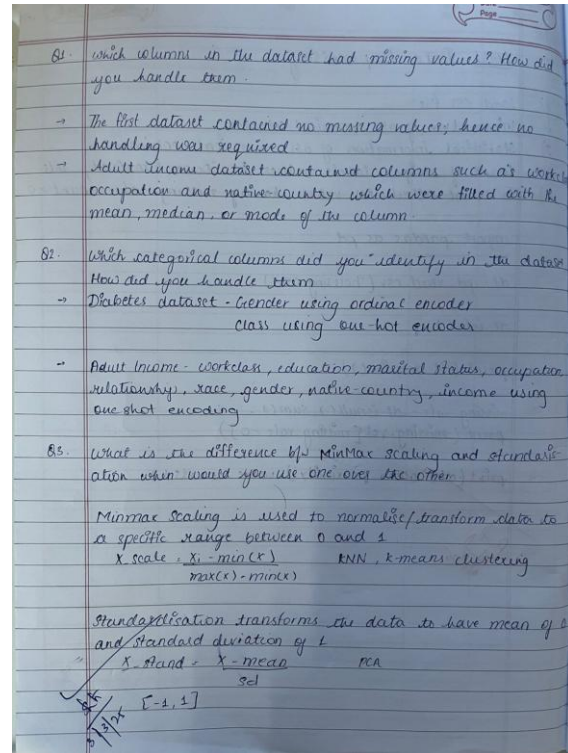
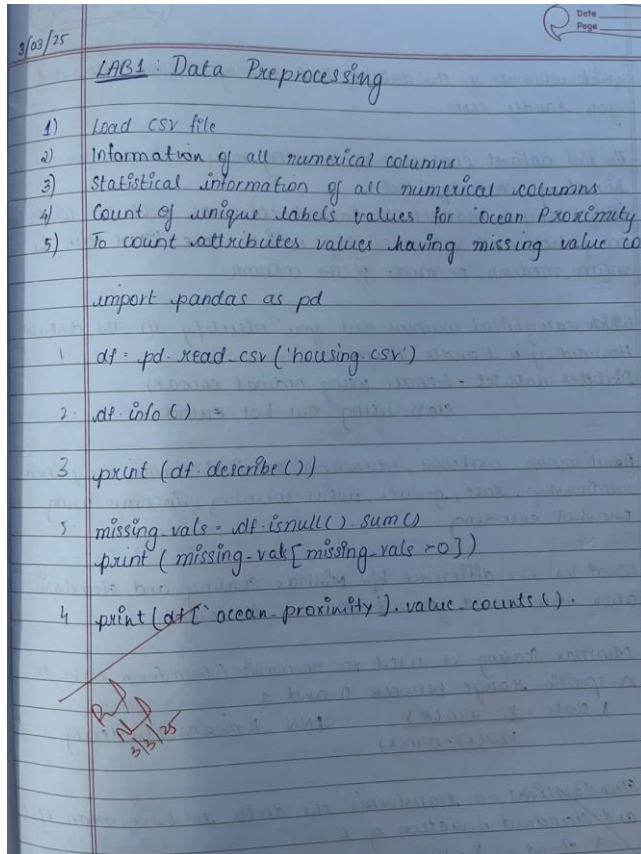
tickers = ["HDFCBANK.NS", "ICICI.NS", "KOTAKBANK.NS"]
data = yf.download(tickers, start="2024-01-01", end="2024-12-30",
group_by='ticker')
print("First 5 rows of the dataset:")
print(data.head())
print("\nShape of the dataset:")
print(data.shape)
print("\nColumn names:")
print(data.columns)
print("\n")
reliance_data = data['KOTAKBANK.NS']
print("\nSummary statistics for Reliance Industries:")
print(reliance_data.describe())
reliance_data['Daily Return'] = reliance_data['Close'].pct_change()
print("\n")
plt.figure(figsize=(12, 6))
plt.subplot(2, 1, 1)
reliance_data['Close'].plot(title="KOTAK Industries - Closing Price")
plt.subplot(2, 1, 2)
reliance_data['Daily Return'].plot(title="kotak Industries - Daily Returns", color='red')
plt.tight_layout()
plt.show()
reliance_data.to_csv('kotak_stock_data.csv')
print("\nkotak stock data saved to 'kotak_stock_data.csv'.")

```

Program 2

Demonstrate various data pre-processing techniques for a given dataset

Screenshot:



Code:

```
from google.colab import files
diabetes=files.upload()
```

```
from google.colab import files
adult_income=files.upload()
```

```
df1=pd.read_csv("Dataset of Diabetes .csv")
df1.head()
```

```
df2=pd.read_csv("adult.csv")
df2.head()
```

```
df1.info()
df2.info()
```



```

df1.describe()
df2.describe()

missing_values1 = df1.isnull().sum()
print(missing_values1)
missing_values2 = df2.isnull().sum()
print(missing_values2)

df1['Gender'] = df1['Gender'].replace('f', 'F')
ordinal_encoder = OrdinalEncoder(categories=[["F", "M"]])
df1["Gender_Encoded"] =
ordinal_encoder.fit_transform(df1[["Gender"]]) onehot_encoder =
OneHotEncoder()
encoded_data =
onehot_encoder.fit_transform(df1[["CLASS"]]) encoded_array
= encoded_data.toarray()
encoded_df = pd.DataFrame(encoded_array,
columns=onehot_encoder.get_feature_names_out(["CLASS"])) df_encoded = pd.concat([df1, encoded_df],
axis=1)
df1 = pd.concat([df1, encoded_df], axis=1)
df1.drop("CLASS", axis=1, inplace=True)
df1.drop("Gender", axis=1, inplace=True)
print(df2.head())
from sklearn.preprocessing import OrdinalEncoder, OneHotEncoder
df_copy2 = df2
ordinal_encoder = OrdinalEncoder(categories=[["Male", "Female"]])
df_copy2["Gender_Encoded"] =
ordinal_encoder.fit_transform(df_copy2[["gender"]])
print(df_copy2[["gender", "Gender_Encoded"]])

onehot_encoder = OneHotEncoder()
encoded_data =
onehot_encoder.fit_transform(df2[["occupation", "workclass", "education", "marital-
status", "relationship", "race", "native-country", "income"]])
encoded_array = encoded_data.toarray()
encoded_df =
pd.DataFrame(encoded_array,
columns=onehot_encoder.get_feature_names_out(["occupation", "workclass", "education", "marital-
status", "relationship", "race", "native-country", "income"]))
df_encoded = pd.concat([df_copy2, encoded_df], axis=1)

df_encoded.drop("gender", axis=1, inplace=True)
df_encoded.drop("occupation", axis=1, inplace=True)
df_encoded.drop("workclass", axis=1, inplace=True)
df_encoded.drop("education", axis=1, inplace=True)
df_encoded.drop("marital-status", axis=1, inplace=True)
df_encoded.drop("relationship", axis=1, inplace=True)
df_encoded.drop("race", axis=1, inplace=True)

```

```

df_encoded.drop("native-country", axis=1, inplace=True)
df_encoded.drop("income", axis=1, inplace=True)
print(df_encoded.head())

normalizer = MinMaxScaler()
df_encoded[["fnlwgt", "educational-num", "capital-gain", "capital-loss", "hours-per-week"]] =
normalizer.fit_transform(df_encoded[["fnlwgt", "educational-num", "capital-gain", "capital-loss", "hours-per-
week"]
])
df_encoded.head()
normalizer = MinMaxScaler()
df1[["No_Pation", "AGE", "Urea", "Cr" , "HbA1c" , "Chol", "TG", "HDL", "LDL", "VLDL", "BMI"]] =
normalizer.fit_transform(df1[["No_Pation", "AGE", "Urea", "Cr" , "HbA1c" ,
"Chol", "TG", "HDL", "LDL", "VLDL", "BMI"]])
df1.head()

```

Program 3

Implement Linear and Multi-Linear Regression algorithm using appropriate dataset

Screenshot:

classmate
Date _____
Page _____

LAB 2:
Solve the linear regression of the data of the work and product sales

x_i work	y_i (sales in thousands)
1	2
2	4
3	5
4	9

$$X = \begin{bmatrix} 1 & 1 \\ 1 & 2 \\ 1 & 3 \\ 1 & 4 \end{bmatrix} \quad Y = \begin{bmatrix} 2 \\ 4 \\ 5 \\ 9 \end{bmatrix}$$

$$\beta = (X^T X)^{-1} X^T Y$$

$$X^T X = \begin{bmatrix} 4 & 10 \\ 10 & 30 \end{bmatrix} \quad (X^T X)^{-1} = \begin{bmatrix} 1.5 & -0.5 \\ -0.5 & 0.2 \end{bmatrix}$$

$$(X^T X)^{-1} X^T = \begin{bmatrix} 1.5 & -0.5 \\ -0.5 & 0.2 \end{bmatrix} \begin{bmatrix} 1 & 1 & 1 & 1 \\ 1 & 2 & 3 & 4 \end{bmatrix}$$

$$= \begin{bmatrix} 1 & 0.5 & 0 & -0.5 \\ -0.3 & -0.1 & 0.1 & 0.3 \end{bmatrix}$$

$$[(X^T X)^{-1} X^T] Y = \begin{bmatrix} 1 & 0.5 & 0 & -0.5 \\ -0.3 & -0.1 & 0.1 & 0.3 \end{bmatrix} \begin{bmatrix} 2 \\ 4 \\ 5 \\ 9 \end{bmatrix} = \begin{bmatrix} -0.5 \\ 2.2 \end{bmatrix} \text{ intercept}$$

~~regression eqⁿ = $y = a_0 + a_1 x$
= $-0.5 + 2.2 x$~~

~~Q1.~~

Q1. Considering the 3 datasets (canada per capita income.csv, hiring.csv, hiring.csv), did you perform any data preprocessing steps (e.g. missing values, scaling etc) if yes, why.

Ans. dataset 1:

- no missing values were present, so no imputation
- Data was already numeric

2. dataset 2:

missing values in 'experience' column were handled using median

dataset 3:

missing values in experience were handled by filling in median
missing values in test score were replaced by mean.

Q2. For canada-per capita income.csv did you visualize the regression line along with the datapoints? What does the plot tell you about the relationship between year and per capita

- There is positive correlation between year and per capita income
- As year increases, per capita income increases, indicating an upward trend.

Q3. For hiring.csv what is the predicted salary for a candidate with 12 years of experience, 10 test scores and 10 interview score?

Solⁿ \$87405.80

Q4. For 1000 companies.csv, did you encode variables (ex. state)? If yes, how

801 Scaling was not done in this model

one-hot was ~~not~~ used

label encoding was done using ~~label encoder~~ () since ML models require numerical values.

9/11
10/3/21
10/10

Code:

```
from google.colab import files
per_capita_income=files.upload()

from google.colab import files
salary=files.upload()

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import OrdinalEncoder, OneHotEncoder
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from scipy import stats
from sklearn import linear_model

df1=pd.read_csv("canada_per_capita_income.csv")
df1.head()

df2=pd.read_csv("salary.csv")
df2.head()
df2.YearsExperience.median()
df2.YearsExperience =
df2.YearsExperience.fillna(df2.YearsExperience.median()) df2

plt.xlabel("year")
plt.ylabel("per capita income (US$)")
plt.scatter(df1.year, df1['per capita income (US$)'])

plt.xlabel("YearsExperience")
plt.ylabel("Salary")
plt.scatter(df2.YearsExperience, df2.Salary)

reg1 = linear_model.LinearRegression()
reg1.intercept_
reg1.predict([[2020]])

reg2 = linear_model.LinearRegression()
reg2.fit(df2.drop('Salary', axis='columns'), df2['Salary'])
reg2.coef_
reg2.intercept_
reg2.predict([[12]])

from google.colab import files
hiring=files.upload()

from google.colab import files
companies=files.upload()

df3=pd.read_csv("hiring.csv")
```

```

df3.head()

df4=pd.read_csv("1000_Companies.csv")
df4.head()

df3.isnull().sum()
df4.isnull().sum()

df3_copy = df3.copy()
experience_mapping = {'two': 2, 'three': 3, 'five': 5, 'seven': 7, 'ten': 10, 'eleven': 11}
df3_copy['experience'] = df3_copy['experience'].map(experience_mapping)
median_experience = df3_copy['experience'].median()
df3_copy['experience'] = df3_copy['experience'].fillna(median_experience)
df3_copy
df3_copy['test_score(out of 10)'] = df3_copy['test_score(out of 10)'].fillna(df3_copy['test_score(out of 10)'].mean())
reg3 = linear_model.LinearRegression()
reg3.fit(df3_copy.drop('salary($)', axis='columns'),
df3_copy['salary($)']) reg3.coef_
reg3.intercept_
reg3.predict([[2,9,6]])
reg3.predict([[12,10,10]])
ohe = OneHotEncoder(sparse_output=False, handle_unknown='ignore') state_encoded =
ohe.fit_transform(df4[['State']])
state_encoded_df = pd.DataFrame(state_encoded, columns=ohe.get_feature_names_out(['State']))

df4 = pd.concat([df4, state_encoded_df], axis=1).drop(columns=['State'])
print(df4)
reg4 = linear_model.LinearRegression()
reg4.fit(df4.drop('Profit',axis='columns'),df4.Profit)
print(reg4.coef_)
print(reg4.intercept_)
reg4.predict([[91694.48, 515841.3, 11931.24,0,1,0]])

```


Program 4

Build Logistic Regression Model for a given dataset

Screenshot:

19/05/24

classmate

Date _____
Page _____

Logistic Regression

Q. Consider a binary classification problem where we want to predict whether a 'student' will pass or fail based on their study hours. The logistic regression model has been trained and the learned parameters are $a_0 = 5$ (intercept) and $a_1 = 0.8$ (coefficient for study hours).

(a) Write the logistic regression equation for this problem.

$$p(x) = \frac{1}{1 + e^{-(a_0 + a_1 x)}}$$

(b) Calculate the probability that 10 students who studied for 7 hours a day will pass.

$$x = 7$$

$$p(x) = p(7) = \frac{1}{1 + e^{-(5 + 0.8(7))}}$$

$$= 0.645$$

(c) Determine the predicted class [pass or fail] based on the threshold 0.5.

$$p(7) = 0.645 \geq 0.5$$

Student will pass

Q2. Consider $x = [2, 1, 0]$ for three classes. Apply softmax function to find the probability values of three classes.

Solⁿ. Softmax: give $x = [2, 1, 0]$

$$\text{softmax}(x_1) = \frac{e^2}{e^2 + e^1 + e^0} = 0.665$$

$$\text{softmax}(x_2) = \frac{e^1}{e^2 + e^1 + e^0} = 0.244$$

$$\text{softmax}(x_3) = \frac{e^0}{e^2 + e^1 + e^0} = 0.091$$

where $\text{softmax}(x) = \frac{e^x}{\sum_{j=1}^J e^{x_j}}$

Q1. For dataset 'HE-lemma-sep.csv'

- Satisfaction level - Employees with low satisfaction levels were more likely to leave.
- Time spent in company - employees who spent more years in the company are more likely to leave.
- Average monthly hours - high/extreme low hours showed higher chance of leaving.
- Number of projects - Employees with too many or too less project are more likely to leave.
- Salary level - Employees with low salary are more likely to leave.
- Departement - Certain departments had higher turnover rates due to role demands.

(ii) What was the accuracy of the logistic regression model? Is the accuracy good? Why or why not?

Solⁿ accuracy = 76.26% which is good considering the size of the dataset. But room for improvement exists.

Q2. For xoo dataset.

- The only data preprocessing that was done was to map the integer datatype in xoo data to that of categorical value in xoo class type.
- No missing data was present.
- The confusion matrix reveals that the model classification almost all the class types correctly, except for a single data point, hence the prediction of model is very accurate (97.1%).
- 'investrator' was being predicted as 'uptile' which is most likely due to the features being similar to each other.

Code:

```
from google.colab import files
hr=files.upload()

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import OrdinalEncoder, OneHotEncoder
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from scipy import stats
from sklearn import linear_model
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

df1=pd.read_csv("HR_comma_sep.csv")
df1.head()
df1.isnull().sum()
plt.figure(figsize=(12, 6))
sns.barplot(x='Department', y='left', data=df1)
plt.title('Employee Retention Rate by Department')
plt.xlabel('Department')
plt.ylabel('Proportion of Employees Left')
plt.xticks(rotation=45, ha='right')
plt.show()

ohe = OneHotEncoder(handle_unknown='ignore', sparse_output=False)
department_encoded = ohe.fit_transform(df1[['Department']])
department_encoded_df = pd.DataFrame(department_encoded,
columns=ohe.get_feature_names_out(['Department']))
df1 = pd.concat([df1, department_encoded_df], axis=1)
df1 = df1.drop('Department', axis=1)
ordinal_encoder = OrdinalEncoder(categories=[['low', 'medium', 'high']], dtype=np.int64)
salary_encoded = ordinal_encoder.fit_transform(df1[['salary']])
df1['salary_encoded'] = salary_encoded
df1 = df1.drop('salary', axis=1)
df1.head()

correlation_matrix = df1.corr()
plt.figure(figsize=(12, 10))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Matrix of Features')
plt.show()
plt.figure(figsize=(8, 6))
sns.barplot(x='salary_encoded', y='left', data=df1)
plt.title('Impact of Employee Salary on Retention')
plt.xlabel('Salary Level (Encoded)')
plt.ylabel('Proportion of Employees Left')
plt.show()
```



```

df_copy = df1[['number_project', 'average_monthly_hours', 'time_spend_company', 'left', 'salary_encoded',
'satisfaction_level', 'Work_accident']]
df_copy.head()
X = df_copy.drop('left', axis=1)
y = df_copy['left']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
model = LogisticRegression(max_iter=1000)
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy of the Logistic Regression model: {accuracy}")

from google.colab import files
zoodata=files.upload()
zootype=files.upload()

zoo_data = pd.read_csv('zoo-data.csv')
zoo_class = pd.read_csv('zoo-class-type.csv')
merged_data = pd.merge(zoo_data, zoo_class, left_on='class_type', right_on='Class_Number')
merged_data = merged_data.drop(['Animal_Names', 'Number_Of_Animal_Species_In_Class',
'Class_Number', 'class_type', 'animal_name'], axis=1)
X = merged_data.drop('Class_Type', axis=1)
y = merged_data['Class_Type']
print(merged_data.head())
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
model = LogisticRegression(max_iter=1000)
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy: {accuracy}")
cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=np.unique(y_test))
disp.plot(cmap="Blues", values_format="d")
plt.title("Confusion Matrix")
plt.show()

```

Program 5

Use an appropriate data set for building the decision tree (ID3) and apply this knowledge to classify a new sample.

Screenshot:

classmate
Date _____
Page _____

24/04/25
LAB 04: Decision Trees

Q. Consider the following dataset, calculate entropy, gain w.r.t target variable 'classification'. Identify whether splitting node should be a_1 , or a_2

Instance	a_1	a_2	classification
1	Hot	High	No
2	Hot	High	No
6	Cool	High	No
7	Hot	High	No
8	Hot	Normal	Yes

Solⁿ: Entropy (S) = $-P_0 \log_2 P_0 - P_1 \log_2 P_1$
 Entropy (S) = $[-\frac{4}{5} \log_2 \frac{4}{5} - \frac{1}{5} \log_2 \frac{1}{5}] = 0.7219$

for a_1
 Entropy (S_H) = $[-\frac{4}{4} \log_2 \frac{4}{4} - \frac{1}{4} \log_2 \frac{1}{4}] = 0.8113$
 (S_L) = $[-\frac{1}{1} \log_2 \frac{1}{1}] = 0$

Information gain = Entropy (S) - $\sum_{\text{regions}} \frac{|S_v|}{|S|} \cdot \text{Entropy}(S_v)$
 = $0.7219 - \left\{ \frac{4}{5} (0.8113) + \frac{1}{5} (0) \right\} = 0.7219 - 0.6420 = 0.0799$

for a_2
 Entropy (S_H) = $[-\frac{4}{4} \log_2 \frac{4}{4}] = 0$
 Entropy (S_L) = $[-\frac{1}{1} \log_2 \frac{1}{1}] = 0$

Information gain = Entropy (S) - $\sum_{\text{regions}} \frac{|S_v|}{|S|} \cdot \text{Entropy}(S_v)$
 = $0.7219 - 0 - 0 = 0.7219$ ✓ max info gain

choose a_2 for splitting variable

classmate
Date _____
Page _____

24/04/25

Q1. For Iris dataset
 what was the accuracy? $\rightarrow 93\%$ [0.93]
 (i) What does the confusion matrix tell about model performance?
 (ii) Were there any misclassifications? Which classes were most confused?

Ans: rows \rightarrow actual class
 columns \rightarrow predicted class

	1	0	0
0	9	1	
0	1	9	

iris-setosa were classified correctly, however 1 versicolour was misclassified as virginica and 1 virginica was also misclassified as versicolour.

Q2. For petrol consumption dataset
 Can you interpret the Regression Tree structure? what are the most important features for predicting petrol consumption? How does the Regression Tree handle continuous target variables compared to the Decision Tree classifier?

Ans: Regression tree is a decision tree used for decision tasks. Hence it predicts continuous values rather than discrete categories.

\rightarrow There are 4 most important features: population-driven license [0.651569], Average income [0.740522], petrol tax [0.065150], paved highways [0.042460]

\rightarrow The Regression Tree keeps splitting until the variance (MSE) is minimised in each region

• Each leaf node, final prediction is the mean [or weighted mean] of all target variables in that region.

Code:

```
from google.colab import files
iris=files.upload()
df1=pd.read_csv("iris.csv")
df1.head()

df1.isnull().sum()

X = df1.drop('species', axis=1)
y = df1['species']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
clf = DecisionTreeClassifier(criterion='entropy')
clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy:.2f}')
print(classification_report(y_test, y_pred))
plt.figure(figsize=(12, 8))
plot_tree(clf, filled=True, feature_names=X.columns,
class_names=y.unique()) plt.show()

cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=clf.classes_)
cmap = plt.cm.get_cmap('PuBuGn')
disp.plot(cmap=cmap)
plt.show()

drug=files.upload()
df2=pd.read_csv("drug.csv")
df2.head()
df2.isnull().sum()

label_encoders = { }
for column in df2.columns:
    le = LabelEncoder()
    df2[column] = le.fit_transform(df2[column])
    label_encoders[column] = le
X = df2.drop('Drug', axis=1)
y = df2['Drug']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
clf = DecisionTreeClassifier(criterion='entropy')
clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f'Accuracy: {accuracy:.2f}')
print(classification_report(y_test, y_pred))
plt.figure(figsize=(12, 8))
plot_tree(clf, filled=True, feature_names=X.columns, class_names=[str(c) for c in y.unique()])
plt.show()

cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=clf.classes_)
```

```

cmap = plt.cm.Blues
disp.plot(cmap=cmap)
plt.show()

pc=files.upload()
df3=pd.read_csv("petrol_consumption.csv")
df3.head()
df3.isnull().sum()
X = df3.drop('Petrol_Consumption', axis=1)
y = df3['Petrol_Consumption']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
regressor = DecisionTreeRegressor(random_state=42)
regressor.fit(X_train, y_train)
y_pred =
regressor.predict(X_test)
mse = mean_squared_error(y_test, y_pred)
rmse = sqrt(mse)
mae = mean_absolute_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
print(f'Mean Squared Error: {mse:.2f}')
print(f'Root Mean Squared Error:
{rmse:.2f}') print(f'Mean Absolute Error:
{mae:.2f}') print(f'R-squared: {r2:.2f}')
plt.figure(figsize=(30, 30))
plot_tree(regressor, filled=True, feature_names=X.columns, fontsize=10)
plt.show()

```

Program 6

Build KNN Classification model for a given dataset.

Screenshot

K-Nearest Neighbours

Consider the following dataset, for $k=3$ and test data $(X, 35, 100)$ as $(Person, Age, Salary)$ solve using knn classifier model and predict the target.

Person	Age	Salary	Target	Distance	Rank
A	18	50	N	52.81	5
B	23	55	N	46.57	4
C	24	70	N	31.95	2
D	71	60	Y	40.44	3
E	43	70	Y	31.04	1
F	38	40	Y	60.07	6
X	35	100	?		

Euclidean Distance : $\sqrt{(x_2 - x_1)^2 + (y_2 - y_1)^2}$

For $k=3$

Rank 1 [E] = Y
Rank 2 [C] = N
Rank 3 [D] = Y

majority [Y]

∴ According to KNN the target for X is Y

For Iris dataset :

How to choose the k value? Demonstrate using accuracy rate and error rate.

Choosing the k -value is generally done by taking the square root of number of entities in the dataset (often the nearest odd number to the square root is taken as ' k ', this is to avoid binary classification).

Optimal ' k ' - where the test/validation accuracy is the highest with accuracy rate.

Optimal ' k ' with error rate, is at the point where error rate is minimum.

Diabetes Dataset

What is the purpose of feature scaling? How to perform it?

Sol: It is essential in KNN, as KNN is a distance based algorithm and without scaling, features with larger ranges would influence the calculations disproportionately. This causes the model to be biased towards features with larger values resulting in poor performance.

→ To overcome this feature scaling is done.

Code:

```
from google.colab import files
iris=files.upload()
df1=pd.read_csv("iris (2).csv")
df1.head()
df1.isnull().sum()
X = df1.drop('species', axis=1)
y = df1['species']
```

```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
best_k = 1
best_accuracy = 0
for k in range(1,
11):
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train, y_train)
    y_pred = knn.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    print(f"Accuracy for k={k}: {accuracy}, Error Rate for k={k}: {1-accuracy}")
    if accuracy > best_accuracy:
        best_accuracy = accuracy
        best_k = k
print(f"Best k value: {best_k}")
knn = KNeighborsClassifier(n_neighbors=3)
knn.fit(X_train, y_train)
y_pred = knn.predict(X_test)
print("Accuracy Score:", accuracy_score(y_test, y_pred))
print("\nConfusion Matrix:")
cm = confusion_matrix(y_test, y_pred)
print(cm)
print("\nClassification Report:")
print(classification_report(y_test, y_pred))
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
            xticklabels=knn.classes_, yticklabels=knn.classes_)
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()

diabetes=files.upload()
df2=pd.read_csv("diabetes.csv")
df2.head()
df2.isnull().sum()
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_scaled = scaler.fit_transform(df2.drop('Outcome', axis=1))
X_train, X_test, y_train, y_test = train_test_split(X_scaled, df2['Outcome'], test_size=0.2, random_state=42)
best_k = 1
best_accuracy = 0
for k in range(1,
11):
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train, y_train)
    y_pred = knn.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    print(f"Accuracy for k={k}: {accuracy}")
    if accuracy > best_accuracy:
        best_accuracy = accuracy
        best_k = k
print(f"Best k value: {best_k}")

```

```

knn = KNeighborsClassifier(n_neighbors=best_k) knn.fit(X_train, y_train)
y_pred = knn.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
cm = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:")
print(cm)
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.xlabel("Predicted") plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
print("\nClassification Report:")
print(classification_report(y_test, y_pred))

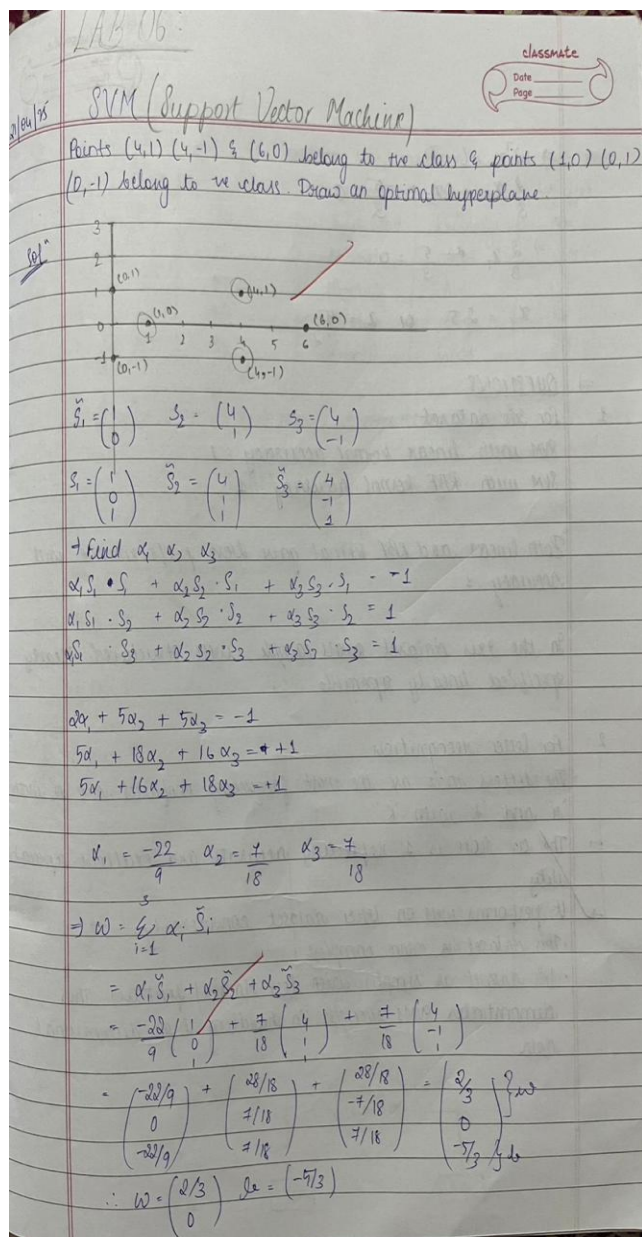
heart=files.upload()
df3=pd.read_csv("heart.csv")
df3.head()
df3.isnull().sum()
X = df3.drop('target', axis=1)
y = df3['target']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
best_k = 1
best_accuracy = 0
for k in range(1,
11):
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(X_train, y_train)
    y_pred = knn.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    print(f"Accuracy for k={k}: {accuracy}, Error Rate for k={k}: {1-accuracy}")
    if accuracy > best_accuracy:
        best_accuracy = accuracy
        best_k = k
print(f"Best k value: {best_k}")
knn = KNeighborsClassifier(n_neighbors=optimal_k)
knn.fit(X_train, y_train)
y_pred = knn.predict(X_test)
print("Accuracy Score:", accuracy_score(y_test, y_pred))
print("\nConfusion Matrix:")
cm = confusion_matrix(y_test, y_pred)
print(cm)
print("\nClassification Report:")
print(classification_report(y_test, y_pred))
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
            xticklabels=knn.classes_, yticklabels=knn.classes_)
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()

```


Program 7

Build Support vector machine model for a given dataset

Screenshot:



$\Rightarrow W^T x + b = 0$

$$\frac{2}{3} x_1 + 0x_2 + \frac{-5}{3} = 0$$

$$\Rightarrow \frac{2}{3} x_1 - \frac{5}{3} = 0$$

$$x_1 = 2.5 \quad \text{or} \quad x_1 = \frac{5}{2}$$

\Rightarrow QUESTIONS:

1. For Iris dataset
 SVM with linear kernel accuracy: 1
 SVM with RBF kernel accuracy: 1

Both linear and RBF kernel gave best performance with accuracy: 1

In the Iris dataset small sample, well structured, clearly specified linearly separable.

2. For Letter-recognition
 The letters that are the most frequently confused are 'p' with 'i' and 'k' with 'r'.

\rightarrow The score is 1 reflecting accurate and excellent separability.

\rightarrow It performs well on letter dataset considerably.

- This dataset is more complex.
- Iris dataset is simpler with less classes/features. This demonstrates SVM's strength in handling high dimensional data.

Code:

```
from google.colab import files
iris=files.upload()
df1=pd.read_csv("iris (1).csv")
df1.head()
X = df1.drop('species', axis=1)
y = df1['species']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
rbf_svm = SVC(kernel='rbf')
rbf_svm.fit(X_train, y_train)
rbf_y_pred = rbf_svm.predict(X_test)
print("RBF Kernel SVM:")
print("Accuracy:", accuracy_score(y_test, rbf_y_pred))
cm = confusion_matrix(y_test, rbf_y_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap="Blues")
plt.title('Confusion Matrix for RBF Kernel SVM')
plt.xlabel('Predicted')
plt.ylabel('True') plt.show()
print(classification_report(y_test, rbf_y_pred))
linear_svm = SVC(kernel='linear')
linear_svm.fit(X_train, y_train)
linear_y_pred = linear_svm.predict(X_test)
print("\nLinear Kernel SVM:")
print("Accuracy:", accuracy_score(y_test, linear_y_pred))
cm = confusion_matrix(y_test, linear_y_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap="Blues")
plt.title('Confusion Matrix for Linear Kernel SVM')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()
print(classification_report(y_test, linear_y_pred))
letter=files.upload()
df2=pd.read_csv("letter-recognition.csv")
df2.head()
X = df2.drop('letter', axis=1)
y = df2['letter']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
svm_classifier = SVC(kernel='linear', probability=True)
svm_classifier.fit(X_train, y_train)
y_pred =
svm_classifier.predict(X_test)
print("Accuracy:", accuracy_score(y_test, y_pred))
print(classification_report(y_test, y_pred))
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(10,10))
sns.heatmap(cm, annot=True, fmt='d', cmap="Blues")
plt.title('Confusion Matrix for SVM')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()
lb = LabelBinarizer()
```

```
lb.fit(y_test)
```

```
y_test_lb = lb.transform(y_test)
y_pred_prob =
svm_classifier.predict_proba(X_test) fpr = { }
tpr = { }
thresh = { }
roc_auc = dict()
n_class = y_test_lb.shape[1]
for i in range(n_class):
    fpr[i], tpr[i], thresh[i] = roc_curve(y_test_lb[:,i], y_pred_prob[:,i])
    roc_auc[i] = auc(fpr[i], tpr[i])
plt.plot(fpr[0], tpr[0], linestyle='--',color='orange', label='SVM (AUC = %0.2f)' % roc_auc[0])
plt.title('ROC Curve for Class 0')
plt.xlabel('False Positive
Rate') plt.ylabel('True Positive
rate') plt.legend(loc='best')
plt.show()
print(f"AUC score for class 0: {roc_auc[0]}")
```

Program 8

Implement Random forest ensemble method on a given dataset

Screenshot:

LAB 07: Random Forest

For Sample S₁, draw the decision tree, consider CGPA as root node.

S.No	CGPA	Interactiveness	Communication Skills	Practical knowledge	Job offer
1	≥ 9	Yes	Good	Good	Yes
2	< 9	No	Moderate	Good	Yes
3	≥ 9	No	Moderate	Avg	No
4	≥ 9	No	Moderate	Avg	No
5	≥ 9	Yes	Moderate	Good	Yes

Considering CGPA as root node, Job offer as target node.

```

graph TD
    CGPA{CGPA} -- "≥ 9" --> Interactiveness{Interactiveness}
    CGPA -- "< 9" --> Yes1[Yes]
    Interactiveness -- "Yes" --> Yes2[Yes]
    Interactiveness -- "No" --> No[No]
  
```

Sample S₂

S.No	CGPA	Interactiveness	Communication Skills	Practical knowledge	Job offer
2	< 9	No	Moderate	Good	Yes
3	≥ 9	No	Moderate	Average	No
3	≥ 9	No	Moderate	Average	No
5	≥ 9	Yes	Moderate	Good	Yes
5	≥ 9	Yes	Moderate	Good	Yes

```

graph TD
    Interactiveness{Interactiveness} -- "Yes" --> Yes3[Yes]
    Interactiveness -- "No" --> Knowledge{Practical knowledge}
    Knowledge -- "Yes" --> Yes4[Yes]
    Knowledge -- "No" --> No2[No]
  
```

QUESTIONS:

- For iris.csv dataset:
- The best accuracy for this dataset is 100%.
- It was found using 9 estimator trees and fine-tuning showed that adding any more trees beyond 9 doesn't affect the model performance.
- The final confusion matrix found was:

19	0	0
0	13	0
0	0	13

Code:

```

from google.colab import files
iris=files.upload()
df1=pd.read_csv("iris (4).csv")
df1.head()
X = df1.drop('species', axis=1)
y = df1['species']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)
rf_classifier = RandomForestClassifier(random_state=0)
rf_classifier.fit(X_train, y_train)
y_pred =
rf_classifier.predict(X_test)
default_accuracy = accuracy_score(y_test, y_pred)
print(f"Accuracy with default n_estimators: {default_accuracy}")
best_accuracy = 0
best_n_estimators = 0
for n_estimators in range(1, 101):
    rf_classifier = RandomForestClassifier(n_estimators=n_estimators, random_state=0)
    rf_classifier.fit(X_train, y_train)
    y_pred = rf_classifier.predict(X_test)
  
```

```
accuracy = accuracy_score(y_test, y_pred)
if accuracy > best_accuracy:best_accuracy
= accuracy best_n_estimators =
n_estimators
print(f"\nBest accuracy: {best_accuracy} achieved with n_estimators = {best_n_estimators}")
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
            xticklabels=np.unique(y_test), yticklabels=np.unique(y_test))
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()
```

Program 9

Implement Boosting ensemble method on a given dataset

Screenshot:

05/06/25 Ada Boost

for the following sample data, show the decision stump calculation steps for the attribute CAPA.

CAPA	Interactiveness	Practical	Communication	Job	weight
>9	Yes	knowledge low	skill (good)	offer	$\frac{1}{6} \times 0.1249$
<9	No	Good	Moderate	Yes	$\frac{1}{6} \times 0.2501$
>9	No	Average	Moderate	No	$\frac{1}{6} \times 0.2501$
<9	No	Average	Good	No	$\frac{1}{6} \times 0.1249$
>9	Yes	Good	Moderate	Yes	$\frac{1}{6} \times 0.1249$
>9	Yes	Good	Moderate	Yes	$\frac{1}{6} \times 0.1249$

for CAPA

Error = $\frac{2 \times 1}{3} = 0.333$

$$\alpha_{CAPA} = \frac{1}{2} \ln \left[\frac{1 - \text{Error}}{\text{Error}} \right]$$

$$= \frac{1}{2} \ln \left(\frac{1 - 0.333}{0.333} \right) = 0.347$$

Normalizing factor $Z_{CAPA} = W_{\text{correct}} \times e^{-\alpha_{CAPA}} + W_{\text{incorrect}} \times e^{\alpha_{CAPA}}$

$$Z_{CAPA} = \frac{1}{6} \times 4 \times e^{-0.347} + \frac{1}{6} \times 2 \times e^{0.347}$$

$$Z_{CAPA} = 0.9428$$

new updated weights

$$w_1(d_j)_{\text{correct}} = \frac{w_1(d_j)_{\text{correct}}}{Z_{CAPA}} \times e^{-\alpha_{CAPA}} = \frac{\frac{1}{6}}{0.9428} \times e^{-0.347} = 0.1249$$

$$w_1(d_j)_{\text{incorrect}} = \frac{w_1(d_j)_{\text{incorrect}}}{Z_{CAPA}} \times e^{\alpha_{CAPA}} = \frac{\frac{1}{6}}{0.9428} \times e^{0.347} = 0.2501$$

QUESTION:

for income.csv dataset

→ The best accuracy score is 83%.

→ Confusion matrix

	10658	457
2023	1521	

→ 160 trees were used to improve performance

→ For $n=10$, accuracy = 89.77%

Confusion matrix

	10722	387
2138	1406	

5/5

Code:

```
from google.colab import files
income=files.upload()
df1=pd.read_csv("income.csv")
df1.head()
X = df1.drop('income_level', axis=1)
y = df1['income_level'] X = pd.get_dummies(X)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
abc = AdaBoostClassifier(n_estimators=10, random_state=42)
abc.fit(X_train, y_train)
y_pred = abc.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print(f"Initial AdaBoost accuracy (10 trees): {accuracy}")
param_grid = {'n_estimators': [50, 100, 150, 200]}
grid_search = GridSearchCV(AdaBoostClassifier(random_state=42), param_grid, cv=5, scoring='accuracy')
grid_search.fit(X_train, y_train)
print(f"Best parameters: {grid_search.best_params_}")
print(f"Best cross-validation score: {grid_search.best_score_}")
best_abc = grid_search.best_estimator_
y_pred_best = best_abc.predict(X_test)
best_accuracy = accuracy_score(y_test,
y_pred_best)
print(f"Accuracy of the best model on the test set: {best_accuracy}")
cm = confusion_matrix(y_test, y_pred_best)
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['<=50K', '>50K'], yticklabels=['<=50K', '>50K'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```


Program 10

Build k-Means algorithm to cluster a set of data stored in a .CSV file

Screenshot:

classmate
Date _____
Page _____

LAB 09

K-MEANS ALGORITHM:

Q. Compute 2 clusters using k-means algorithm for clustering where initial cluster centres are (1.0, 1.0) and (5.0, 7.0). Execute for 2 iterations

Record no.	A	B
R ₁	1.0	1.0
R ₂	1.5	2.0
R ₃	3.0	4.0
R ₄	5.0	7.0
R ₅	3.5	5.0
R ₆	4.5	5.0
R ₇	3.5	4.5

Sol: No. of clusters k = 2, centroid for cluster C₁ = (1.0, 1.0) and centroid for C₂ = (5.0, 7.0)

Iteration 1:

Record no.	close to C ₁ (1.0, 1.0)	close to C ₂ (5.0, 7.0)	Assign to cluster
R ₁ (1.1)	dist(R ₁ , C ₁) = 0.0	dist(R ₁ , C ₂) = 7.21	C ₁
R ₂ (1.5, 2)	dist(R ₂ , C ₁) = 1.12	dist(R ₂ , C ₂) = 6.12	C ₁
R ₃ (3, 4)	dist(R ₃ , C ₁) = 3.61	dist(R ₃ , C ₂) = 3.61	C ₁
R ₄ (5, 7)	dist(R ₄ , C ₁) = 7.21	dist(R ₄ , C ₂) = 0.0	C ₂
R ₅ (3.5, 5)	dist(R ₅ , C ₁) = 4.12	dist(R ₅ , C ₂) = 2.5	C ₂
R ₆ (4.5, 5)	dist(R ₆ , C ₁) = 5.31	dist(R ₆ , C ₂) = 2.06	C ₂
R ₇ (3.5, 4.5)	dist(R ₇ , C ₁) = 4.30	dist(R ₇ , C ₂) = 2.92	C ₂

Cluster 1 { R₁, R₂, R₃ } Cluster 2 { R₄, R₅, R₆, R₇ }

C₁ = $\frac{1.0+1.5+3.0}{3}, \frac{1+2+4}{3}$ C₂ = $\frac{5+3.5+3.4}{4}, \frac{7+5+5+4.5}{4}$

= 1.83, 2.33 = 4.12, 5.37

Page _____

Iteration 2

Record no.	close to C ₁ (1.83, 2.33)	close to C ₂ (4.12, 5.37)	Assign to cluster
R ₁ (1, 1)	1.57	5.87	C ₁
R ₂ (1.5, 2)	0.47	4.27	C ₁
R ₃ (3, 4)	2.04	1.77	C ₂
R ₄ (5, 7)	5.64	1.85	C ₂
R ₅ (3.5, 5)	3.15	0.72	C ₂
R ₆ (4.5, 5)	3.78	0.73	C ₂
R ₇ (3.5, 4.5)	2.74	1.07	C ₂

Cluster 1 { R₁, R₂ } Cluster 2 { R₃, R₄, R₅, R₆, R₇ }

C₁ = $\frac{1.0+1.5}{2}, \frac{1.0+2.0}{2}$ C₂ = $\frac{3+5+3.5+4.5+3.5}{5}, \frac{4+7+5+5+4.5}{5}$

= 1.25, 1.5 = 3.9, 5.1

optimal number of clusters = 3

Code:

```
from google.colab import files
iris=files.upload()
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from scipy import stats
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler

df1=pd.read_csv("iris (4).csv")
df1.head()
df = df1.drop(['sepal_length','sepal_width','species'],axis=1)
scaler = StandardScaler()
scaled_df = scaler.fit_transform(df) wcss = []
for i in range(1, 11):
    kmeans = KMeans(n_clusters=i, init='k-means++', max_iter=300, n_init=10, random_state=0)
    kmeans.fit(scaled_df)
    wcss.append(kmeans.inertia_)
plt.plot(range(1, 11), wcss)
plt.title('Elbow Method')
plt.xlabel('Number of clusters')
plt.ylabel('WCSS')
plt.show()
kmeans = KMeans(n_clusters=3, init='k-means++', max_iter=300, n_init=10, random_state=0)
pred_y = kmeans.fit_predict(scaled_df)
df['cluster'] = pred_y
plt.scatter(df['petal_length'], df['petal_width'], c=df['cluster'])
plt.title('Clusters of Iris Flowers')
plt.xlabel('Petal Length')
plt.ylabel('Petal Width')
plt.show()
```


Program 11

Implement Dimensionality reduction using Principal Component Analysis (PCA) method.

Screenshot:

LAB 10:

classmate
Date _____
Page _____

PCA

Feature	Ex:1	Ex:2	Ex:3	Ex:4
x_1	4	8	13	7
x_2	11	4	5	14

$\lambda_1 = 30.3846$ $e_1 = \begin{bmatrix} 0.5574 \\ -0.8303 \end{bmatrix}$ $e_2 = \begin{bmatrix} 0.8303 \\ 0.5574 \end{bmatrix}$

$\lambda_2 = 6.6151$

Solⁿ $\bar{x}_1 = \frac{1}{4} (4+8+13+7) = 8$

$\bar{x}_2 = \frac{1}{4} (11+4+5+14) = 8.5$

Step 2: Covariance matrix

$$S = \begin{bmatrix} \text{cov}(x_1, x_1) & \text{cov}(x_1, x_2) \\ \text{cov}(x_2, x_1) & \text{cov}(x_2, x_2) \end{bmatrix}$$

$$\text{cov}(x_1, x_1) = \frac{1}{N-1} \sum_{k=1}^N (x_{1k} - \bar{x}_1)(x_{1k} - \bar{x}_1)$$

$$= \frac{1}{3} [(4-8)^2 + (8-8)^2 + (13-8)^2 + (7-8)^2]$$

$$= 14$$

$$\text{cov}(x_1, x_2) = \frac{1}{3} [(4-8)(11-8.5) + (8-8)(4-8.5) + (13-8)(5-8.5) + (7-8)(14-8.5)]$$

$$= -11$$

$$\text{cov}(x_2, x_1) = -11$$

$$\text{cov}(x_2, x_2) = 23$$

$$S = \begin{bmatrix} 14 & -11 \\ -11 & 23 \end{bmatrix}$$

$\lambda_1 = 30.3846$ $\lambda_2 = 6.6151$

$u = \begin{bmatrix} u_1 \\ u_2 \end{bmatrix}$ $\begin{bmatrix} 0 \\ 0 \end{bmatrix} = (S - \lambda I) u$

$$= \begin{bmatrix} 14-\lambda & -11 \\ -11 & 23-\lambda \end{bmatrix} \begin{bmatrix} u_1 \\ u_2 \end{bmatrix} = \begin{bmatrix} (14-\lambda)u_1 - 11u_2 \\ -11u_1 + (23-\lambda)u_2 \end{bmatrix}$$

$(14-\lambda)u_1 - 11u_2 = 0$

$-11u_1 + (23-\lambda)u_2 = 0$

Date _____
Page _____

$u = \begin{bmatrix} u_1 \\ u_2 \end{bmatrix} \cdot \frac{u_1}{11} = \frac{u_2}{14-\lambda}$ $u = \begin{bmatrix} 11 \\ 14-\lambda \end{bmatrix}$

$e_1 = \begin{bmatrix} 0.5574 \\ -0.8303 \end{bmatrix}$ $e_2 = \begin{bmatrix} 0.8303 \\ 0.5574 \end{bmatrix}$

$e^T = \begin{bmatrix} \bar{x}_{1k} - \bar{x}_1 \\ \bar{x}_{2k} - \bar{x}_2 \end{bmatrix} = \begin{bmatrix} 0.5574 & -0.8303 \end{bmatrix} \begin{bmatrix} x_{11} - \bar{x}_1 \\ x_{21} - \bar{x}_2 \end{bmatrix}$

$= 0.5574(4-8) - 0.8303(11-8.5)$

$= -4.30535$

Components

Features	Ex:1	Ex:2	Ex:3	Ex:4
x_1	4	8	13	7
x_2	11	4	5	14
FPC	-4.3052	3.281	5.6428	-5.1238

Accuracy Score	Accuracy Score (before PCA)	Accuracy score (After PCA)
Logistic Regression	0.85	0.83
SUM	0.88	0.86
Random Forest	0.90	0.88

13/5

Code:

```
from google.colab import files
heart=files.upload()

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from scipy import stats
import seaborn as sns
from sklearn.preprocessing import LabelEncoder, OneHotEncoder
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
from sklearn.preprocessing import StandardScaler
from sklearn.svm import SVC
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.decomposition import PCA

df1=pd.read_csv("heart (1).csv")
df1.head()
text_cols = df1.select_dtypes(include=['object']).columns
label_encoder = LabelEncoder()
for col in text_cols:
    df1[col] =
label_encoder.fit_transform(df1[col])
print(df1.head())
X = df1.drop('HeartDisease', axis=1)
y = df1['HeartDisease']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
scaler = StandardScaler()
X_train =
scaler.fit_transform(X_train) X_test =
scaler.transform(X_test)
# Support Vector Machine
svm_model = SVC(kernel='linear', random_state=42)
svm_model.fit(X_train, y_train)
svm_predictions = svm_model.predict(X_test)
svm_accuracy = accuracy_score(y_test, svm_predictions)
print(f"SVM Accuracy: {svm_accuracy}")

# Logistic Regression
lr_model = LogisticRegression(random_state=42)
lr_model.fit(X_train, y_train) lr_predictions = lr_model.predict(X_test)
lr_accuracy = accuracy_score(y_test, lr_predictions)
print(f"Logistic Regression Accuracy: {lr_accuracy}")

# Random Forest
rf_model = RandomForestClassifier(random_state=42)
rf_model.fit(X_train, y_train)
rf_predictions = rf_model.predict(X_test)
```

```
rf_accuracy = accuracy_score(y_test, rf_predictions)
print(f"Random Forest Accuracy: {rf_accuracy}")
```

```
models = {
    "SVM": svm_accuracy,
    "Logistic Regression":
        lr_accuracy, "Random Forest":
        rf_accuracy
}
```

```
best_model = max(models, key=models.get)
print(f"\nBest Model: {best_model} with accuracy {models[best_model]}")
pca = PCA(n_components=0.95)
X_train_pca = pca.fit_transform(X_train)
X_test_pca = pca.transform(X_test)
```

```
svm_model_pca = SVC(kernel='linear', random_state=42)
svm_model_pca.fit(X_train_pca, y_train)
svm_predictions_pca = svm_model_pca.predict(X_test_pca)
svm_accuracy_pca = accuracy_score(y_test, svm_predictions_pca)
print(f"SVM Accuracy (with PCA): {svm_accuracy_pca}")
```

```
lr_model_pca = LogisticRegression(random_state=42)
lr_model_pca.fit(X_train_pca, y_train)
lr_predictions_pca = lr_model_pca.predict(X_test_pca)
lr_accuracy_pca = accuracy_score(y_test, lr_predictions_pca)
print(f"Logistic Regression Accuracy (with PCA): {lr_accuracy_pca}")
```

```
rf_model_pca = RandomForestClassifier(random_state=42)
rf_model_pca.fit(X_train_pca, y_train)
rf_predictions_pca = rf_model_pca.predict(X_test_pca)
rf_accuracy_pca = accuracy_score(y_test, rf_predictions_pca)
print(f"Random Forest Accuracy (with PCA): {rf_accuracy_pca}")
```

```
models_pca = {
    "SVM": svm_accuracy_pca,
    "Logistic Regression": lr_accuracy_pca,
    "Random Forest": rf_accuracy_pca
}
```

```
best_model_pca = max(models_pca, key=models_pca.get)
print(f"\nBest Model (with PCA): {best_model_pca} with accuracy {models_pca[best_model_pca]}")
```